

# **Namespaces, labels and annotations**

## What are they and why do we need them?

- **Namespaces** allow you to subdivide a single physical cluster into isolated sections, each of which is called a namespace.
- **Labels:** are key/value pairs that help for grouping objects, and attaching identifying information.
- **Annotations:** are also key/value pairs that hold non-identifying information which can be useful for other tools and libraries.

# Namespaces

- The namespace provides the scope for names. Names of resources within one namespace need to be unique.
- Another reason is to better monitor resources by team.
- **Example:** development, QA, prod namespaces.

## Namespaces (cont.)

By default, Kubernetes starts with the following three namespaces:

- **Default:** holds the default set of pods, services, and deployments used by the cluster.
- **Kube-public:** Namespace for resources that are publicly available/readable by all.
- **Kube-system:** Namespace for objects/resources created by Kubernetes systems.

```
kubectl get namespaces
```

```
kubectl get namespaces --show-labels
```

## How do we create new namespaces?

- Get the current context

```
kubectl config view --minify
```

## How do we create new namespaces? (cont.)

- Set up new namespaces

```
kubectl config set-context dev --namespace=development \  
  --cluster=YOUR_CLUSTER \  
  --user=YOUR_USER
```

```
kubectl config set-context prod --namespace=production \  
  --cluster=YOUR_CLUSTER\  
  --user=YOUR_USER
```

## How do we create new namespaces? (cont.)

- Switch context

```
kubectl config use-context dev
```

- Check that we are using the right context

```
kubectl config current-context
```

- To delete:

- Exit to another context.

- ```
kubectl config delete-context dev
```

# Common use cases for namespaces

- Roles and Responsibilities.
  - **Architect:** sets up namespace strategy, avoiding snowflake namespaces.
  - **Admin:** implement the namespaces strategy defined by the architect.
- Antipatterns:
  - Disordered mushrooming of clusters.
  - One large cluster that cannot be broken down.



## Common use cases for namespaces (cont.)

- Partitioning landscapes: dev vs. test vs. prod
  - Consistent naming across environments.
  - Better resource control.
- Antipatterns:
  - Over-naming: if no staging is done, then no need of a staging namespace.
  - All projects land on a `dev` namespace. There are no hierarchies in namespaces, so it is better to have `projectABC-dev` , `projectABC-prod` , etc.

## Common use cases for namespaces (cont.)

- Customer partitioning for non-multi-tenant scenarios.
  - Separate namespaces for clients/projects.
- Warning:
  - No official mechanism to enforce access controls across namespaces, even if you can hack your way through [using service accounts](#).

## When to avoid namespaces

- **No ability to enforce partitioning across namespaces.** Users may be able to access other resources in the cluster, regardless of their namespace.
- For billing/geographical/compliance constraints, the best way to enforce separation is to set up different clusters (and enforce suitable controls there).
- Namespaces are not hierarchies, and should not be used for versioning.

# Resource quotas

- A resource quota limits the amount of resources that namespace can use.
- Resource quotas can limit anything in the namespace"
  - total count of each type of object
  - the total storage used
  - total memory or CPU usage of containers in the namespace.
- Good practice to set up **resources.requests** and **resources.limits** in the YAML file.

# Demo

Resource quotas

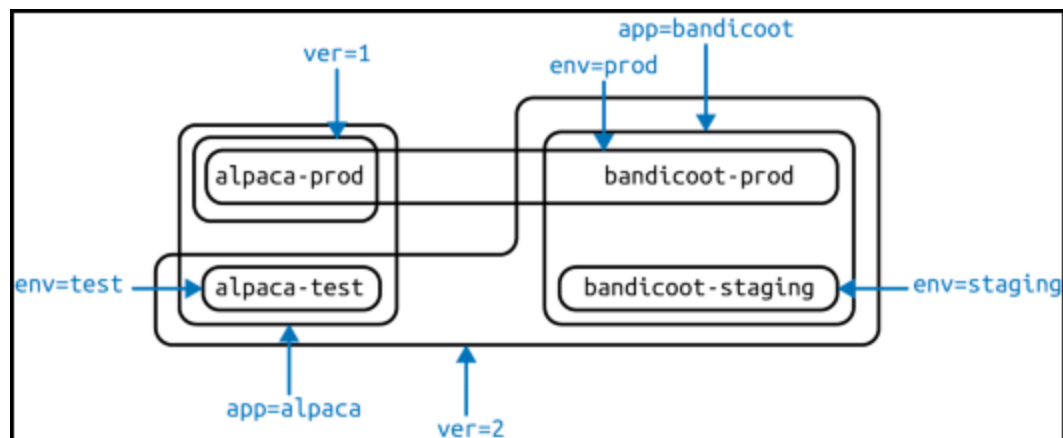
# Labels

## Why labels?

- **Production environment hates singletons:** you may start with a single instance, but then you need to multiply.
- **Hierarchies are not eternal:** Strict hierarchies may complicate scalability. Labels are flexible enough to adapt to these situations.

# Labels syntax: Example

- Let's create two apps, with two environments each.



- Hint:

```
kubectl run alpaca-prod --image=nginx --labels  
="ver=1,app=alpaca,env=prod" --generator=run-pod/v1
```

- Admire our masterpiece:

- `kubectl get pods --show-labels`



## Modifying or updating our labels

- `kubectl label pods alpaca-test "canary=true"`
- `kubectl get pods -L canary`
- `kubectl label pods alpaca-test "canary-"`

## Label selectors

- `kubectl get pods --selector="ver=2"`
- `kubectl get pods --selector="app=bandicoot,ver=2"`
- `kubectl get pods --selector="app in (alpaca,bandicoot)"`

# Selector operators

| Operator                   | Description                        |
|----------------------------|------------------------------------|
| key=value                  | key is set to value                |
| key!=value                 | key is not set to value            |
| key in (value1, value2)    | key is one of value1 or value2     |
| key notin (value1, value2) | key is not one of value1 or value2 |
| key                        | key is set                         |
| !key                       | key is not set                     |

# Annotations

- Add extra metadata:
  - Build, release or image information (Git hashes, timestamp, etc).
  - Keep track of why an object update happened.
  - Way to communicate with other tools (e.g. Ingress).
- Free-form string field.

```
metadata:  
  annotations:  
    icon-url:"https://example.com/icon.png"
```

# ConfigMaps and Secrets

# ConfigMaps

- Small filesystem.
- Set of environment variables/command line in containers.
- Combined with the pod right before, which makes the pod definition itself fully reusable in other environments, by just changing the ConfigMap.
- Can be created in an imperative way or in a declarative way by means of a manifest file.

## Example (Imperative)

```
# my-config.txt sample config file  
parameter1 = value1  
parameter2 = value2
```

```
kubectl create configmap my-config \  
--from-file = my-config.txt \  
--from-literal = extra-param = extra-value \  
--from-literal = another-param = another-value
```

```
# Equivalent YAML file  
kubectl get configmaps my-config -o yaml
```

## Using a ConfigMap

- **Filesystem:** You can mount a ConfigMap into a pod. A file is created for each entry based on the key name. The contents of the file are set to the value.
- **Environment variable/command line:** Dynamically create the command line for a container.



## Example (Filesystem)

```
apiVersion: v1
kind: Pod
metadata:
  name: kuard-config
spec:
  containers:
    - name: test-container
      image: gcr.io/kuar-demo/kuard-amd64:blue
      volumeMounts:
        - name: config-volume
          mountPath: /config
  volumes:
    - name: config-volume
      configMap:
        name: my-config
```

# Example (Environment)

```
apiVersion: v1
kind: Pod
metadata:
  name: kuard-config
spec:
  containers:
    - name: test-container
      image: gcr.io/kuar-demo/kuard-amd64:blue
      command:
        - "/kuard"
        - "$(EXTRA_PARAM)"
      env:
        - name: ANOTHER_PARAM
          valueFrom:
            configMapKeyRef:
              name: my-config
              key: another-param
        - name: EXTRA_PARAM
          valueFrom:
            configMapKeyRef:
              name: my-config
              key: extra-param
```

# Secrets

- Handle extra-sensitive data (passwords, security tokens, private keys).
- By default, K8s secrets are stored as plain text in the `etcd` storage of the cluster. So *anyone with cluster admin rights* can read everything.
- Recent versions have support for encrypting secrets with user-supplied keys (usually integrated into a cloud key store). This allows to skip Kubernetes secrets entirely and rely on the cloud provider's key store.
- **Slides from Google.**

# Demo

Managing configuration maps and secrets.